

# **影像處理、電腦視覺及深度學習概論**

## **(Introduction to Image Processing, Computer Vision and Deep Learning)**

### **Homework 2**

**TA:**

**少鈞: nckubot65904@gmail.com**

**Office Hour: 14:00~16:00, Mon.**

**10:00~12:00, Fri.**

**At CSIE 9F Robotics Lab.**

# Notice (1/2)

- Copying homework is strictly prohibited!! **Penalty: Score will be zero for both persons!!**
- Due date => **09:00:00, 2022/12/12 (Mon.)**

No delay. Penalties for late homework:

- Up to 7 days late, loss of 50% of the score awarded
- After 7 days, the score will be marked as 0.

- You must **attend the demonstration**, or your score will be 0. The demonstration schedule will be announced on NCKU moodle.
- You must **make a GUI** and **follow the format**, or you will get some penalties.
- Upload to => **140.116.154.1 -> Upload/Homework/Hw2**  
**Upload/Homework/Hw2\_05**

➤ User ID: **opencvdl2022**      Password: **opencvdl2022**

## Format

- Filename: Hw2\_StudentID\_Name\_Version.rar  
Hw2\_05\_StudentID\_Name\_Version.rar
  - Ex: Hw2\_F71234567\_林小明\_V1.rar
  - If you want to update your file, you should update your version to be V2.  
Ex: Hw2\_F71234567\_林小明\_V2.rar
- Content: **project folder\*** ( excluding the pictures, only source code )  
\*note: remove your “Debug” folder to reduce file size

# Notice (2/2)

- Python
  - Python 3.7 (<https://www.python.org/downloads/>)
  - opencv-contrib-python (3.4.2.17)
  - Matplotlib 3.1.1
  - UI framework: PyQt5 (5.15.1)

# Assignment scoring (Total: 100%)

## 1. (20%) Image Processing (出題 : Mei)

1.1 (15%) Draw Contour

1.2 (5%) Count Rings

## 2. (20%) Camera Calibration (出題 : Jessica)

2.1 (4%) Corner detection

2.2 (4%) Find the intrinsic matrix

2.3 (4%) Find the extrinsic matrix

2.4 (4%) Find the distortion matrix

2.5 (4%) Show the undistorted result

## 3. (20%) Augmented Reality (出題 : Ming)

3.1 (10%) Show words on board

3.2 (10%) Show words vertically

## 4. (20%) Stereo Disparity Map (出題 : Maton)

4.1 (10%) Stereo Disparity Map

4.2 (10%) Checking the Disparity Value

## 5. (20%) Train a Cat-Dog Classifier Using ResNet50 (出題 : Benjamin)

5.1 (3%) Load the dataset and resize images

5.2 (3%) Plot class distribution of training dataset

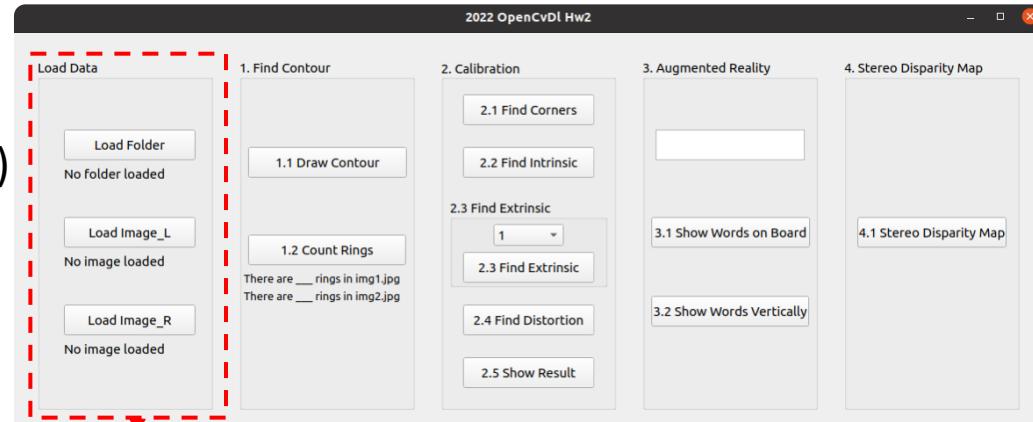
5.3 (3%) Show the structure of ResNet50 model

5.4 (3%) Set up 2 kinds of loss functions to train 2 ResNet50 models

5.5 (3%) Compare the accuracies of 2 ResNet50 models on validation dataset

5.6 (4%) Use the better-trained model to run inference and show the predicted class label

Question 5 needs to upload separately.



Don't fix your data path  
(There is another dataset for demonstration)

# **1. (20%) Find Contour**

(出題 : Mei)

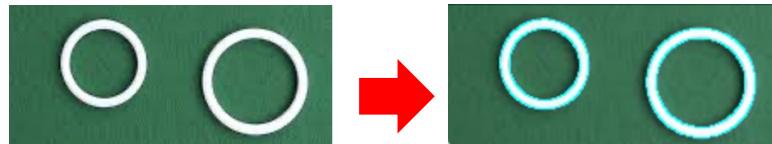
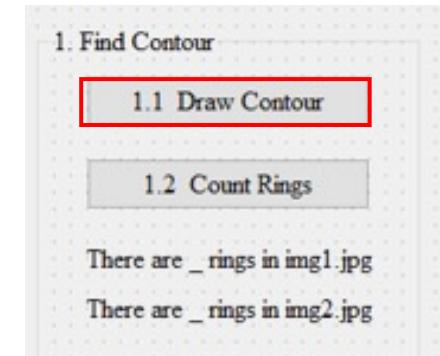
**1.1 (15%) Draw Contour**

**1.2 (5%) Count Rings**

# 1.1 Find Contour – Draw Contour

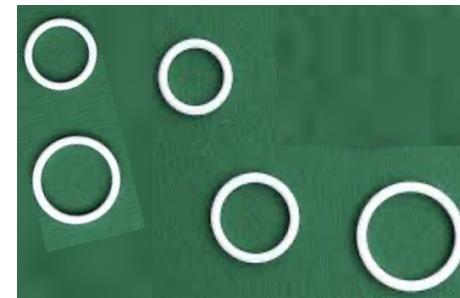
(出題 : Mei)

- Given: two color images, “img1.jpg” and “img2.jpg”
- Q: 1) **Draw Contour**: Using OpenCV functions to find the contours of rings in two images.
- Hint: Textbook Chapter 8, p.234 ~ p.241
  1. RGB → Resize(1/2) → Grayscale
  2. Remember to remove the noise. (use **Gaussian Blur**)
  3. Using some **edge detection functions** to get better results. (cv2.Canny)

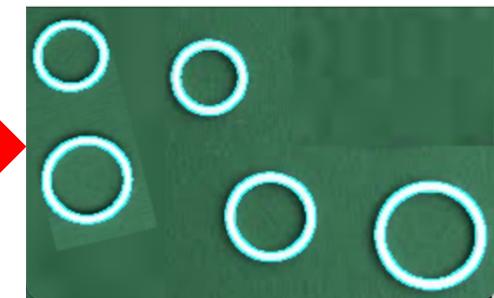


img1.jpg

Draw Contours



img2.jpg

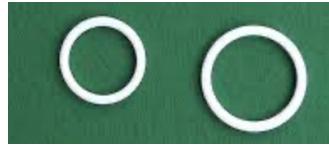


Draw Contours

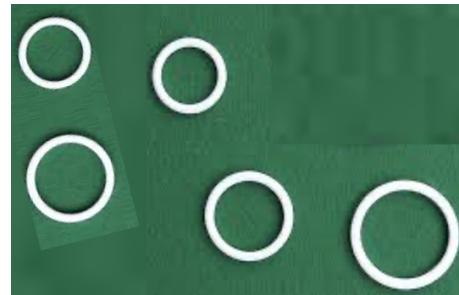
# 1.2 Find Contour – Count Rings

(出題 : Mei)

- Given: two color images, “img1.jpg” and “img2.jpg”
- Q: 2) **Count Rings**: Using OpenCV functions to find how many rings in two images.
- Hint: Textbook Chapter 8, p.234 ~ p.241  
Calculate how many rings



img1.jpg



img2.jpg

1. Find Contour

1.1 Draw Contour

1.2 Count Rings

There are \_ rings in img1.jpg

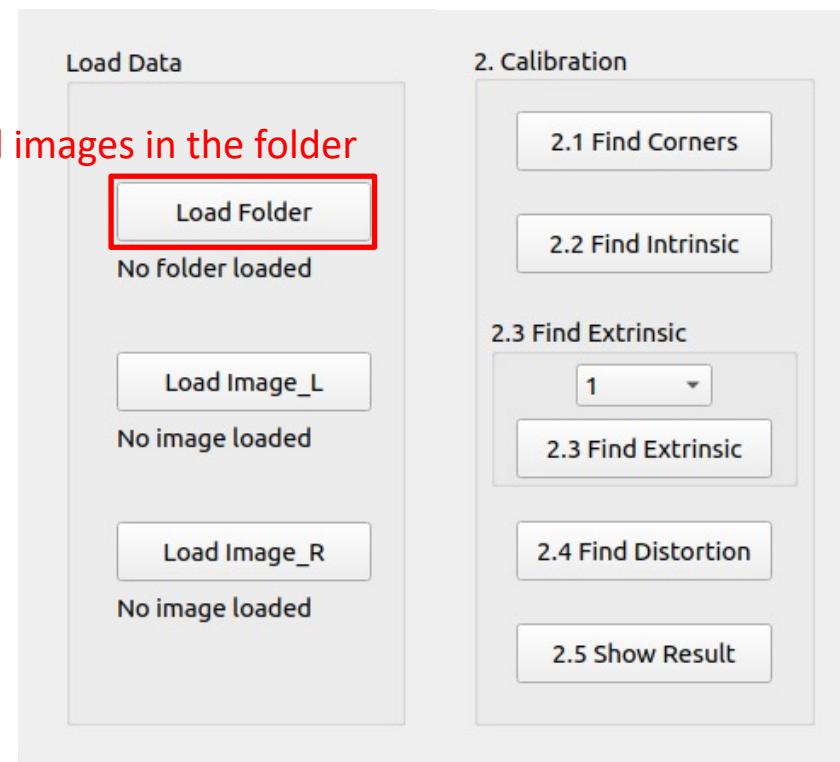
There are \_ rings in img2.jpg

A screenshot of a software interface titled "1. Find Contour". It has three main sections: "1.1 Draw Contour", "1.2 Count Rings" (which is highlighted with a red border), and two text fields below them. The first text field contains the placeholder "There are \_ rings in img1.jpg" and the second contains "There are \_ rings in img2.jpg".

## 2. (20%) Camera Calibration

(出題 : Jessica)

- 2.1 (4%) Corner detection
- 2.2 (4%) Find the intrinsic matrix
- 2.3 (4%) Find the extrinsic matrix
- 2.4 (4%) Find the distortion matrix
- 2.5 (4%) Show the undistorted result



## 2.1 Corner Detection (4%)

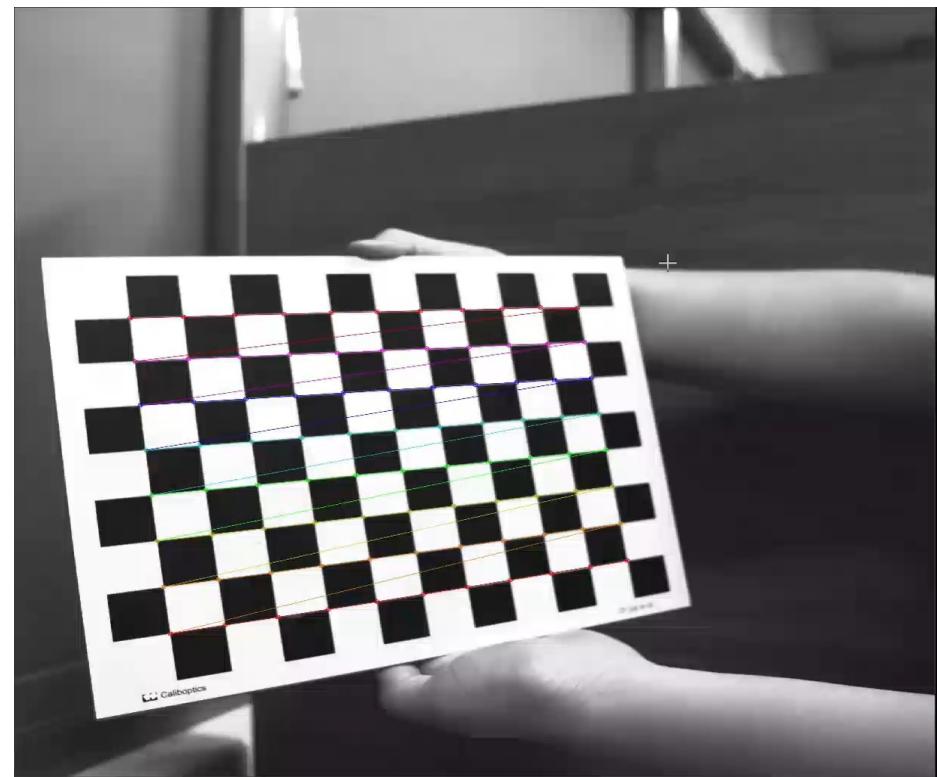
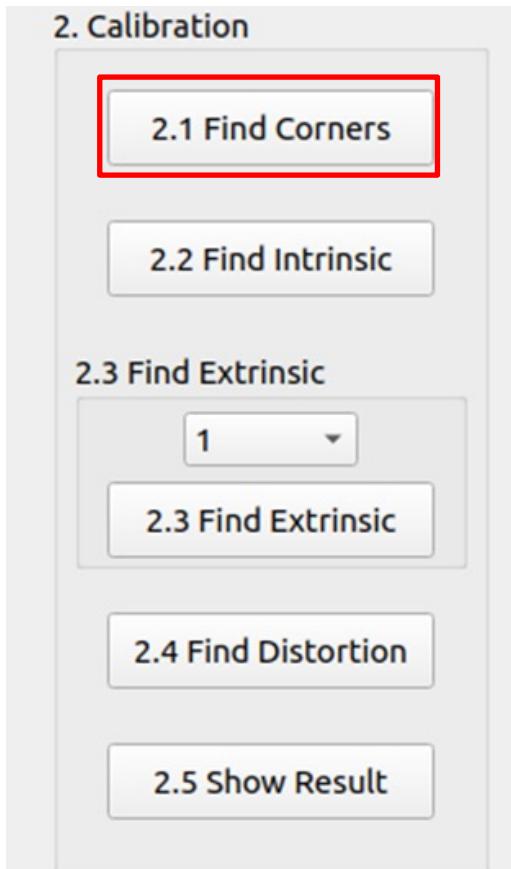
(出題 : Jessica)

- Given: 15 images, 1.bmp ~ 15.bmp
- Q: 1) Find and draw the corners on the chessboard for each image.  
2) Click button “2.1” to show each picture 0.5 seconds.

- Hint :

OpenCV Textbook Chapter 11 (p. 398 ~ p. 399)  
`cv.findChessboardCorners(...)`

- Ex:



## 2.2 Find the Intrinsic Matrix (4%)

(出題 : Jessica)

Given: 15 images, 1.bmp ~ 15.bmp

Q: 1) Find the intrinsic matrix ():  
$$\begin{bmatrix} \alpha & \gamma & u_0 \\ 0 & \beta & v_0 \\ 0 & 0 & 1 \end{bmatrix}$$

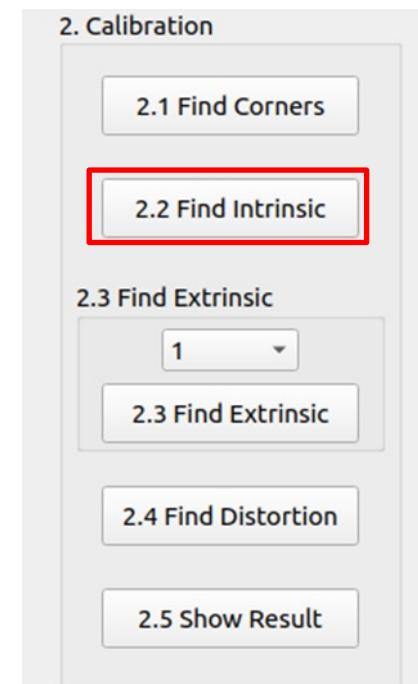
2) Click button “2.2” and then show the result on the console window.

Output format:

```
Intrinsic:  
[[2.22370244e+03 0.00000000e+00 1.03021663e+03]  
 [0.00000000e+00 2.22296836e+03 1.03752624e+03]  
 [0.00000000e+00 0.00000000e+00 1.00000000e+00]]
```

(Just an example)

Hint: OpenCV Textbook Chapter 11 (P.398 ~ p.400)



## 2.3 Find the Extrinsic Matrix (4%)

(出題 : Jessica)

- Given: Intrinsic parameters, distortion coefficients, and the list of 15 images
- Q: 1) Find the extrinsic matrix of the chessboard for each of the 15 images, respectively:

$$\begin{bmatrix} R_{11} & R_{12} & R_{13} & T_1 \\ R_{21} & R_{22} & R_{23} & T_2 \\ R_{31} & R_{32} & R_{33} & T_3 \end{bmatrix}$$

- 2) Click button “2.3” and then show the result on the console window.

Extrinsic:

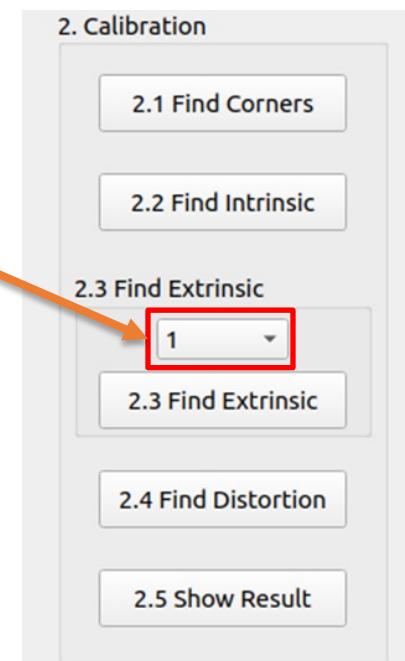
```
[[ -0.8767247 -0.23001438 0.4224301 4.39838495]
 [ 0.19727469 -0.97293475 -0.12033563 0.68022105]
 [ 0.43867585 -0.02216645 0.89837194 16.22126 ]]
```

(Just an example)

- Output format:

- Hint: OpenCV Textbook Chapter 11, p.370~402

- (1) List of numbers: 1~15
- (2) Select 1, then 1.bmp will be applied, and so on



## 2.4 Find the Distortion Matrix (4%)

(出題 : Jessica)

- Given: 15 images
- Q: 1) Find the distortion matrix:  $[k_1, k_2, p_1, p_2, k_3]$   
2) Click button “2.4” to show the result on the console window.

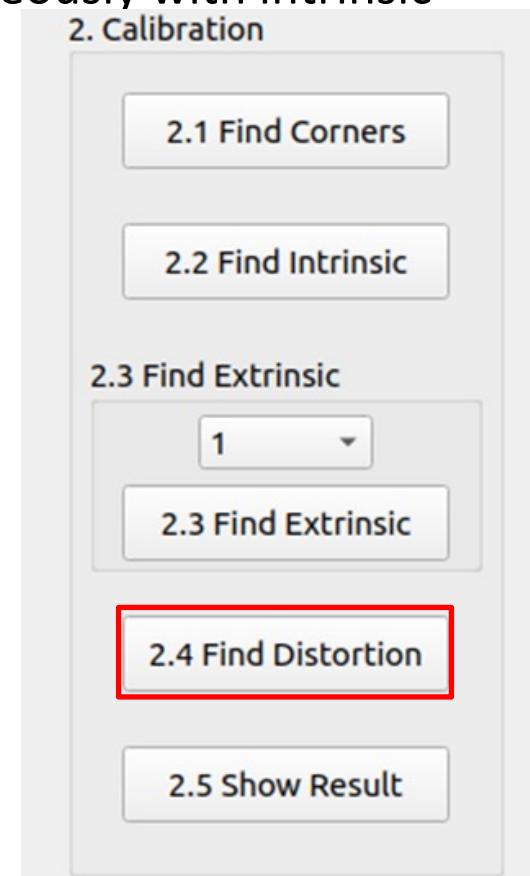
- Output format:

```
Distortion:  
[[ -0.11868112  0.02776881 -0.00092036  0.00047227  0.11793646]]
```

(Just  
an example)

- Hint:

- Distortion coefficients can be obtained simultaneously with intrinsic parameters
  - OpenCV Textbook Chapter 11 (P.398 ~ p.400)



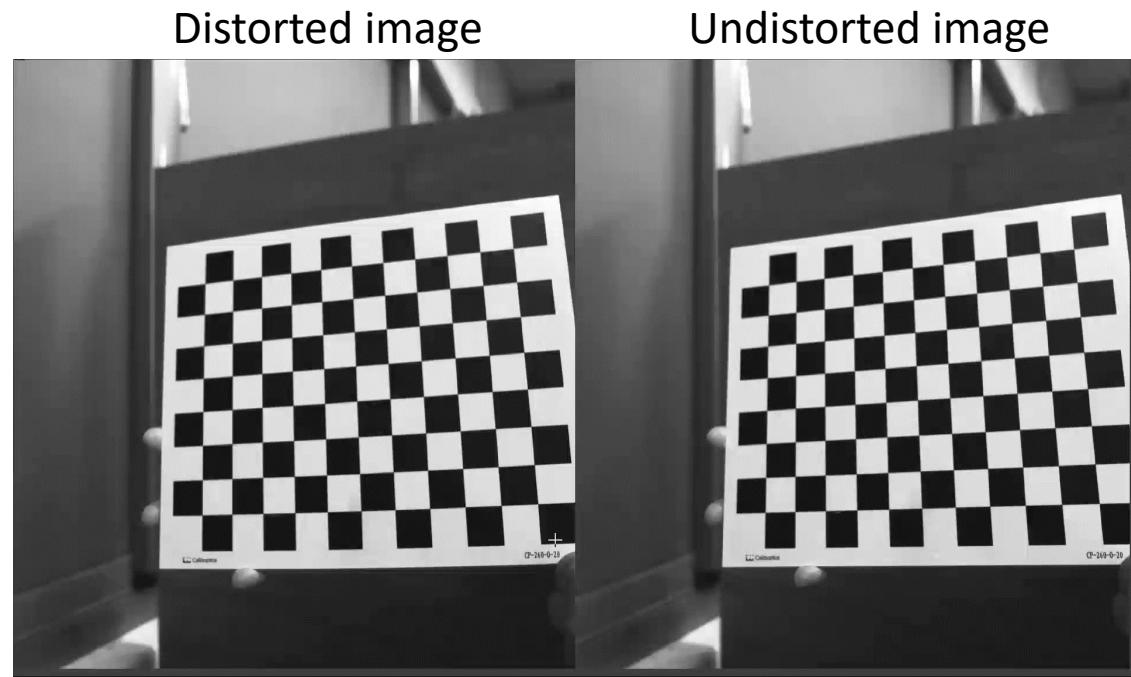
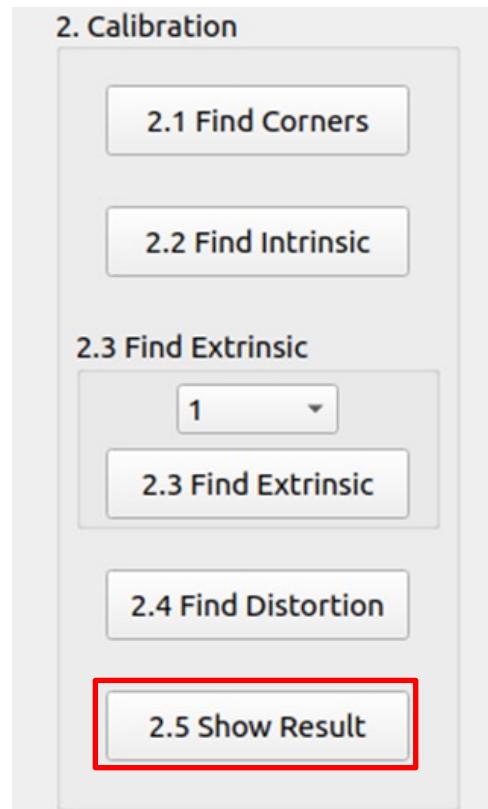
## 2.5 Show the undistorted result (4%)

(出題 : Jessica)

- Given: 15 images
- Q: 1) Undistort the chessboard images.  
2) Show each distorted and undistorted images 0.5 seconds.

- Hint:

- cv::undistort(...) or cv::initUndistortRectifyMap(...)
  - OpenCV Textbook Chapter 11 (P.398 ~ p.400)

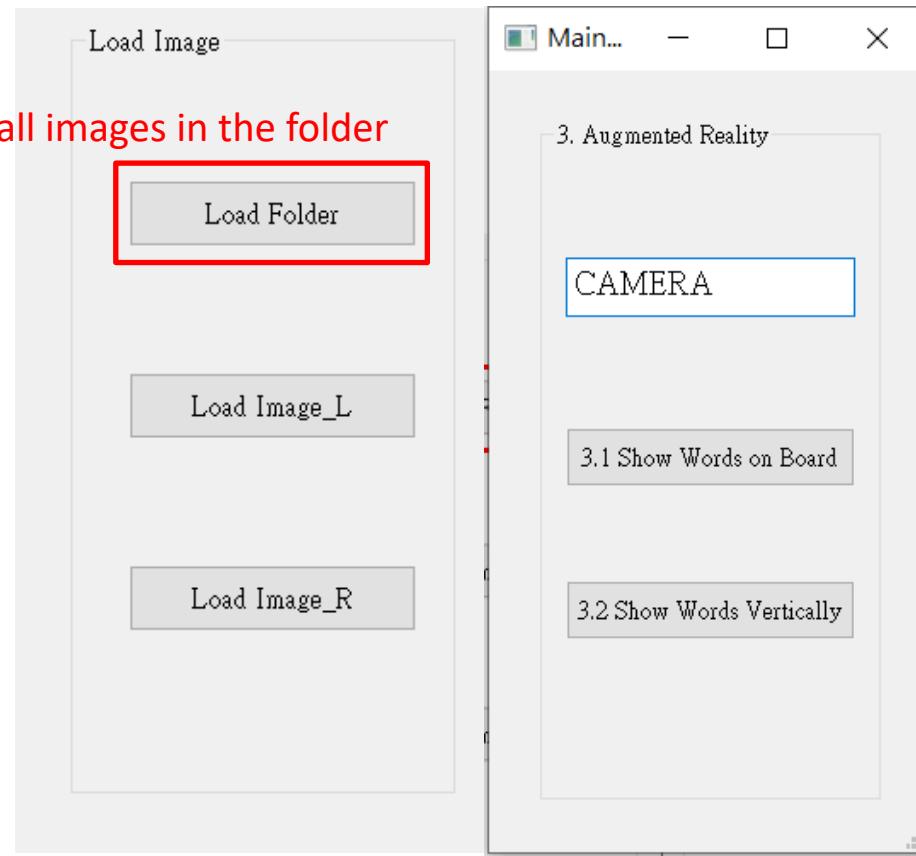


### 3. (20%) Augmented Reality

(出題 : Ming)

3.1 (10%) Show words on board

3.2 (10%) Show words vertically



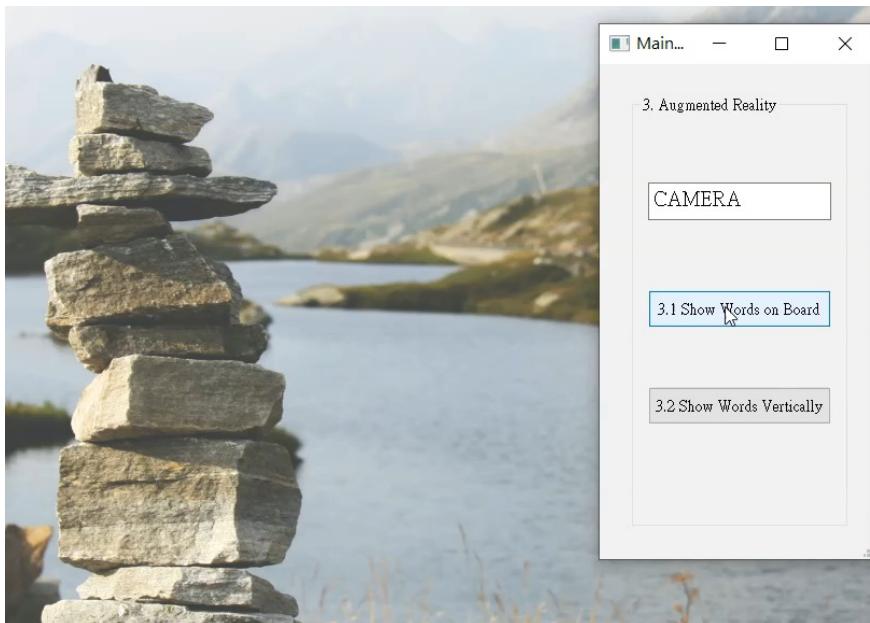
### 3. (20%) Augmented Reality

(出題 : Ming)

- Given: 5 images: 1~5.bmp
- Q:

- 1) Calibrate 5 images to get intrinsic, distortion and extrinsic parameters
- 2) Input a “Word” less than 6 char in English in the textEdit box
- 3) Derive the shape of the “Word” by using the provided library
- 4) Show the “Word” on the chessboards images(1.bmp to 5.bmp)
- 5) Show the “Word” vertically on the chessboards images(1.bmp to 5.bmp)
- 6) Click the button to show the “Word” on the picture. Show each picture for 1 second (total 5 images)

Demo:



Hint : Textbook Chapter 11,  
p.387~395 Calibration  
p.405~412 Projection  
`cv2.calibrateCamera()`  
`cv2.projectPoints()`

# 3. (20%) Augmented Reality

(出題 : Ming)

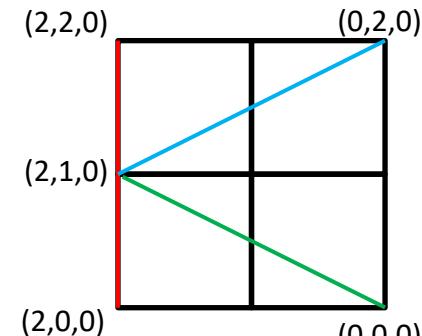
## Guides and Requirements:

- 1) How to use the library: (alphabet\_lib\_onboard.txt, alphabet\_lib\_vertical.txt)
  - Use OpenCV function to read and derive the array or matrix of the char  
Here take 'K' in 'alphabet\_lib\_onboard.txt' for example

Ex (Python):

```
fs = cv2.FileStorage('alphabet_lib_onboard.txt', cv2.FILE_STORAGE_READ)
ch = fs.getNode('K').mat() ➔ get the lines of 'K'
```

```
ch = [[[2, 2, 0], [2, 0, 0]],
       [[0, 2, 0], [2, 1, 0]],
       [[2, 1, 0], [0, 0, 0]]]
```

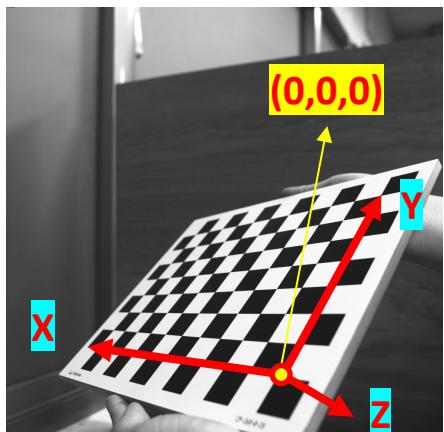


The coordinates are  
in Word Frame

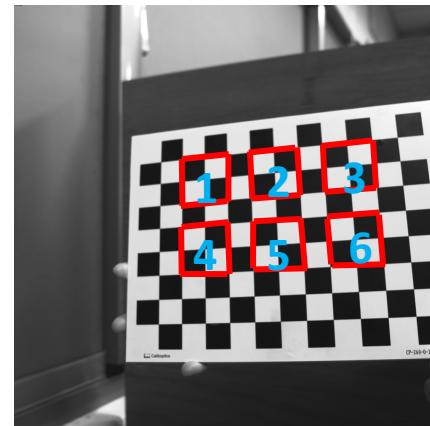
- 'K' consist of 3 lines, so the 'ch array' consists 3 pairs of 3D coordinates in Word Frame representing two ends of the line shown in the upper right image.

## 2) Chessboard Coordinates

- The chessboard x, y, z axis and (0,0,0) coordinate are shown in the bottom left image
- Each Char should be place in the order and position shown in the bottom right image



Chessboard Frame



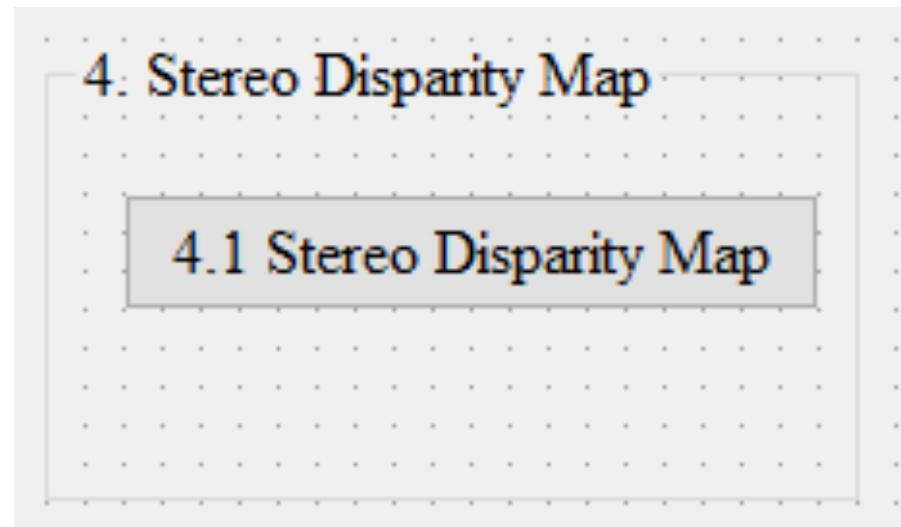
Position and Order

# 4. (20%) Stereo Disparity Map

(出題 : Maton)

## 4.1 (10%) Stereo Disparity Map

## 4.2 (10%) Checking the Disparity Value



# 4.1 (10%) Stereo Disparity Map

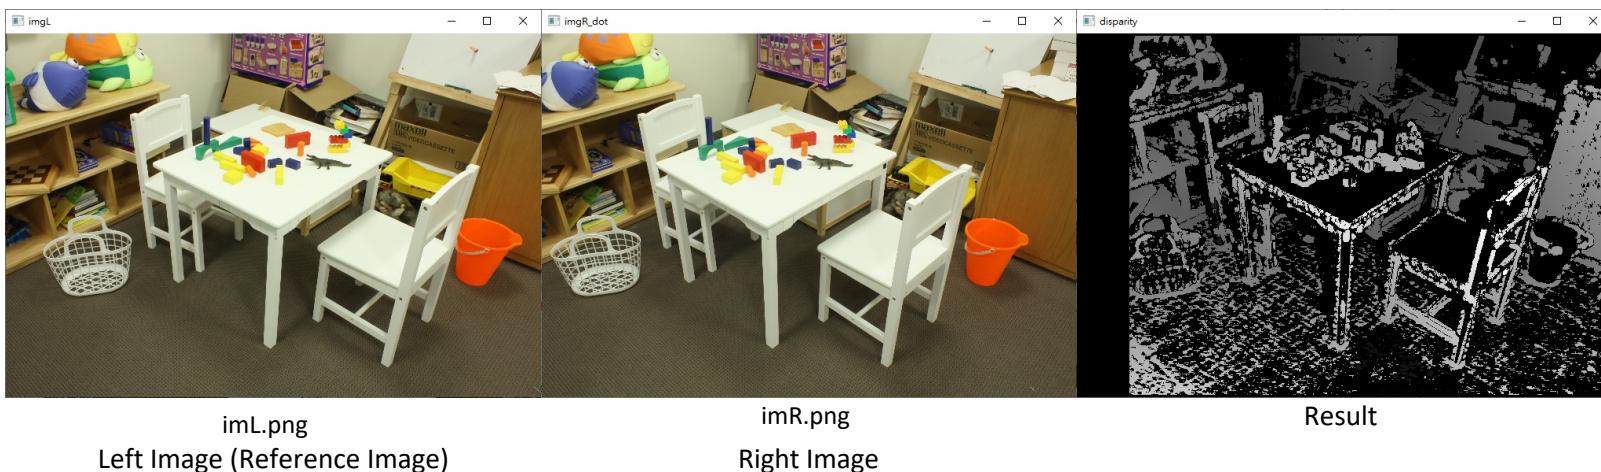
(出題 : Maton)

- Given: a pair of images, imL.png and imR.png (have been rectified)
- Q:
  - Find **the disparity map/image** based on Left and Right stereo images
- Guides:
  - (1) Window Size: Must be **odd** and within the range [5, 255]
  - (2) Search range and direction:
    - Disparity range:
      - Must be **positive** and **divisible by 16**.
      - Map **disparity range** to **gray value range** 0~255 for the purpose of visualization.
    - If the **left image** is the **reference image** (the one used to cal. depth info for each pixel of that Img), then **the search direction at right image** will go **from the right to left** direction.

Camera information: 1) baseline=342.789mm,  
2) focal length=4019.284 pixel,  
3)  $c_x^{right} - c_x^{left} = 279.184$  pixel

OpenCV Textbook Chapter 11 (P.372-373) & OpenCV Textbook Chapter 12 (P.436)

➤ Hint: OpenCV Textbook Chapter 12  
(P.451)  
StereoBM::create(256, 25)



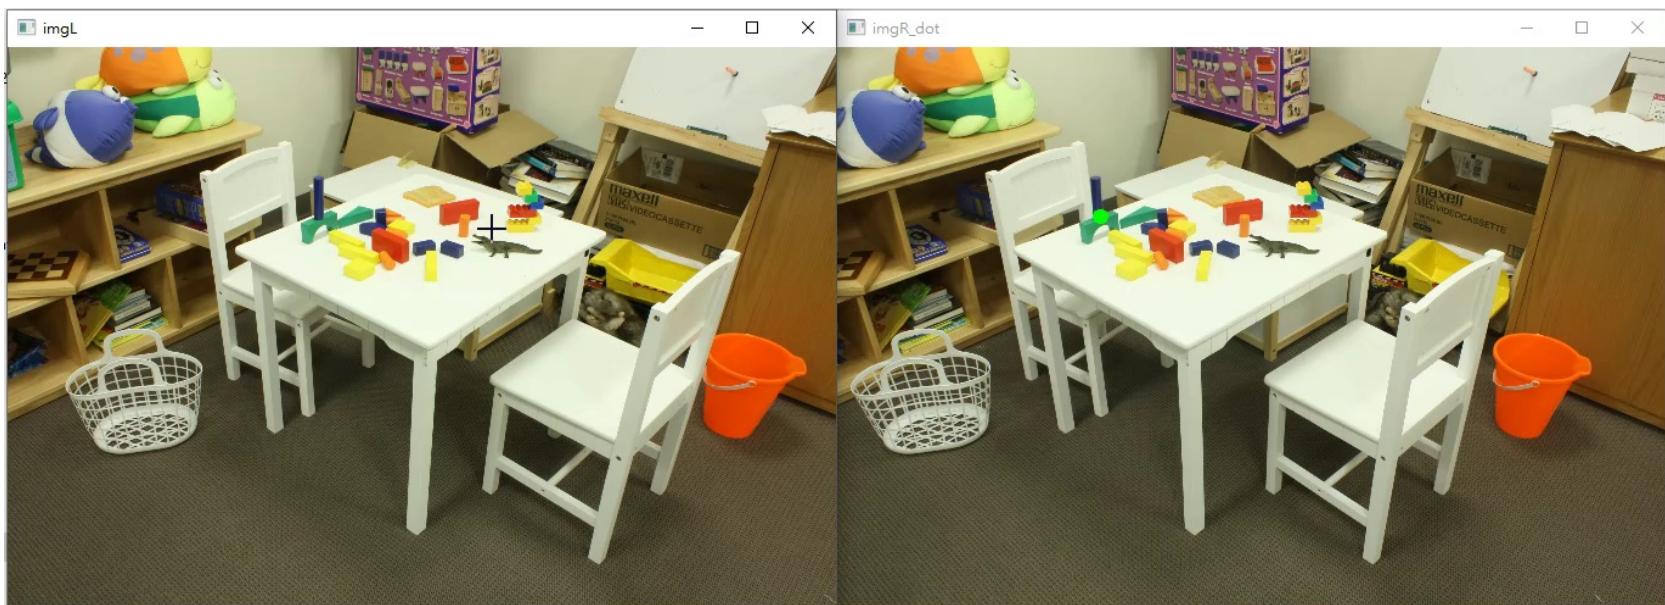
## 4.2 (10%) Checking the Disparity Value

(出題 : Maton)

- Given: a pair of images, imL.png and imR.png and disparity map from Q4.1.
- Q:
  - Click at left image and draw the corresponding dot at right image.
- 1) Click at left image and draw the dot on the right image at accurate position.
- 2) User should allow to repeat 1).

➤ Note: Click at gray position at disparity map result from Q.4.1, ignore the position with 0 disparity(e.g. Failure case).

➤ Result Video:



## 5. Train a Cat-Dog Classifier Using ResNet50 (20%) (出題 : Benjamin)

- 5.1 (3%) Load the dataset and resize images
- 5.2 (3%) Plot class distribution of training dataset
- 5.3 (3%) Show the structure of ResNet50 model
- 5.4 (3%) Set up 2 kinds of loss functions to train 2 ResNet50 models
- 5.5 (4%) Compare the accuracies of 2 ResNet50 models on validation dataset
- 5.6 (4%) Use the better-trained model to run inference and show the predicted class label

# 5.0 Train a Cat-Dog Classifier Using ResNet50

(出題 : Benjamin)

## 1. Objective

- 1) Learn how to train a ResNet50 model to **classify images** of cats and dogs
- 2) Learn how to deal with **imbalanced data**

## 2. ResNet50

- 1) Residual learning: avoid **degradation** problems
- 2) Bottleneck: build a **deeper** network without additional parameters

## 3. Cats and Dogs Dataset (modified from Kaggle Cats and Dogs Dataset)

- 1) Data type: JPG images
- 2) 2 classes: Cat and Dog
- 3) Datasets
  - (1) Path: /Download/Homework/Hw2/Dataset\_OpenCvDI\_Hw2\_Q5.zip
  - (2) Training dataset: 16,200 JPG images in total. However, the dataset is **imbalanced**.
  - (3) Validation dataset: 1,800 JPG images in total.
  - (4) Inference dataset: 10 JPG images in total.  
It is for **testing the inference function** in your GUI program.

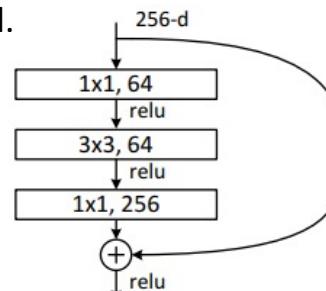
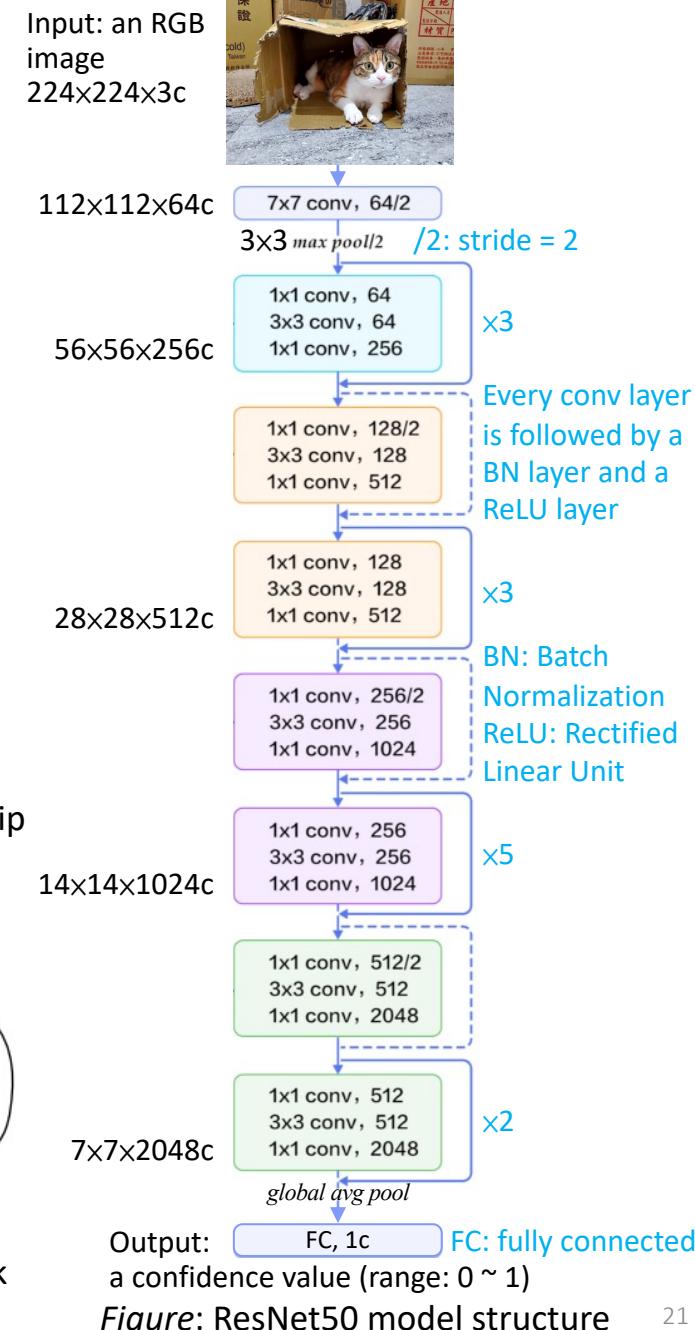


Figure: Residual Block  
(with bottleneck)



## R. Reference

- 1) [Deep Residual Learning for Image Recognition](#)
- 2) [Kaggle Cats and Dogs Dataset](#)

# 5.0 Train a Cat-Dog Classifier Using ResNet50

(出題 : Benjamin)

## 4. Requirements

- 1) Train ResNet50 models with TensorFlow or PyTorch
- 2) Every chart should have a **chart title** and **axis titles**
- 3) Organize the files in this structure:

```
Hw2_05_StudentID_Name_Version // project folder
|-- model          // folder to put trained models
|-- inference_dataset
    |-- Cat
    |-- Dog
|-- main.py        // codes for your GUI program
|-- train.py       // codes for model training
|-- :              // other files or folders you need
```

**Notice: Please include the inference dataset in your homework file.**

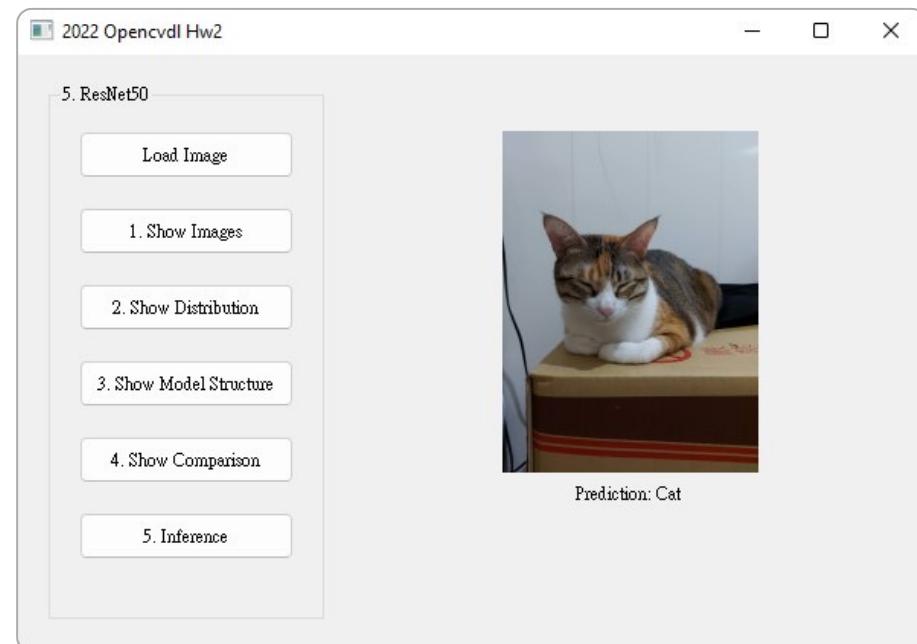


Figure: GUI example

# 5.1 (3%) Load the dataset and resize images

(出題 : Benjamin)

## 1. Objective

### 1) At home:

#### (1) Load the training dataset and validation dataset

→ Hint:

(a) TensorFlow ([tutorial](#)): `tf.keras.utils.image_dataset_from_directory()`

(b) PyTorch ([tutorial](#)): `torch.utils.data.Dataset`

#### (2) Resize images to $224 \times 224 \times 3c$ (RGB)

→ Hint:

(a) TensorFlow: `tf.keras.utils.image_dataset_from_directory(image_size=(?, ?))`

(b) PyTorch ([tutorial](#)): `torchvision.transforms`

### 2) When the demo:

#### (1) Click the button “1. Show Images”

#### (2) Load the inference dataset

#### (3) Resize images to $224 \times 224 \times 3c$ (RGB)

#### (4) Get 1 image from each class in the inference dataset

#### (5) Show images in a new window

→ Hint: use `matplotlib.pyplot` functions to show

images ([tutorial](#)):

(a) `figure()`

(b) `imshow()`

(c) `subplot()`

(d) `title()`

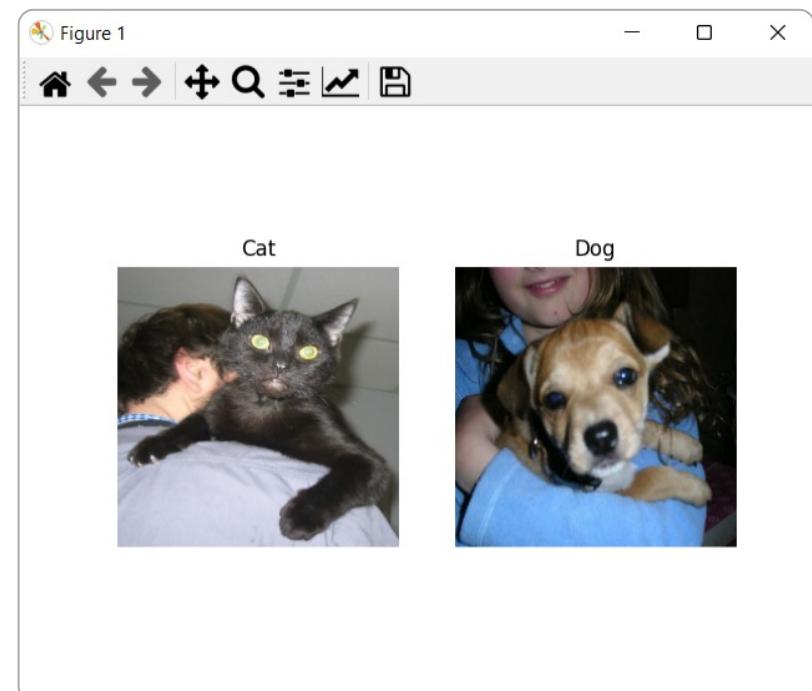
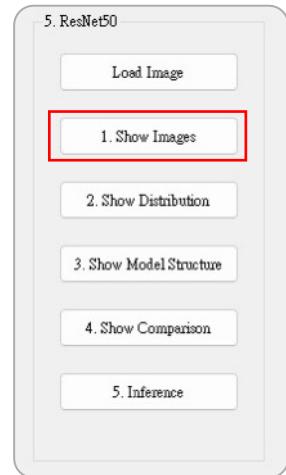


Figure: 1 image from each class

Notice: this is an example, the images might differ

# 5.2 (3%) Plot class distribution of training dataset (出題 : Benjamin)

## 1. Objective

1) At home:

- (1) Load the training dataset
- (2) Iterate through the dataset and count **the number of images of each class**
- (3) Plot the class distribution

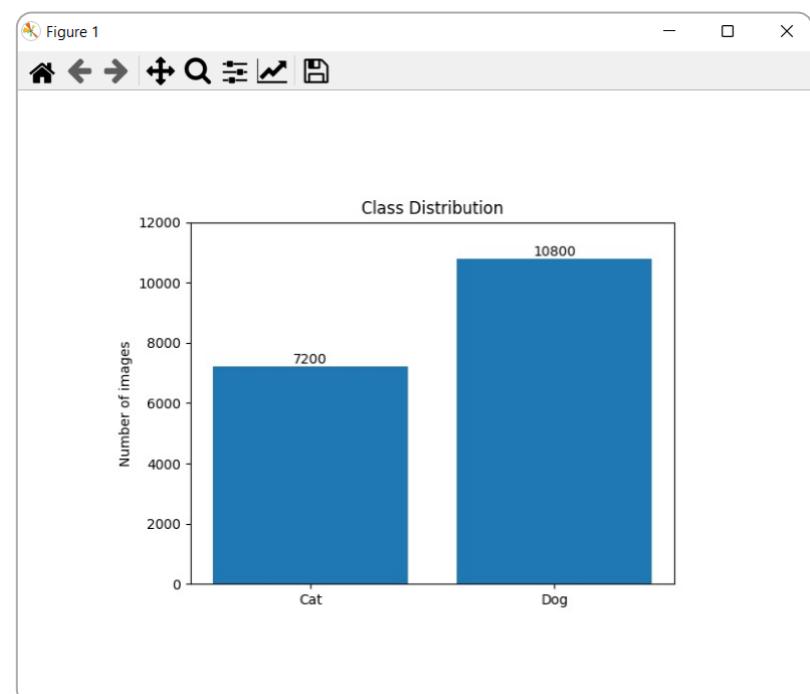
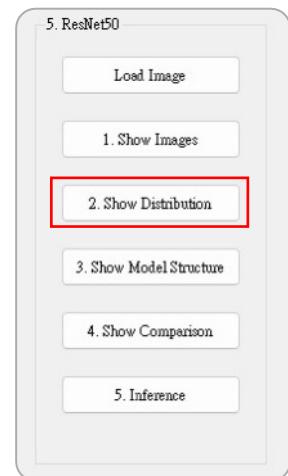
→ Hint: use matplotlib.pyplot functions to plot

- (a) figure()
- (b) bar()
- (c) xticks()
- (d) ylabel()

(4) **Save the figure**

2) When the demo:

- (1) Click the button "**2. Show Distribution**"
- (2) Show the **saved figure** of class distribution in a new window



*Figure: Class Distribution*

Notice: this is an example, the numbers might differ

# 5.3 (3%) Show the structure of ResNet50 model

(出題 : Benjamin)

## 1. Objective

### 1) At home:

#### (1) Build a ResNet50 model

→ Hint:

(a) TensorFlow: `tf.keras.applications.resnet50.ResNet50()`

(b) PyTorch: `torchvision.models.resnet50()`

#### (2) Replace the output layer to a FC (Fully Connected) layer of **1 node** with a **Sigmoid** activation function

→ Hint:

(a) TensorFlow ([tutorial](#)): `tf.keras.layers.Dense(1, activation='sigmoid')`

(b) PyTorch ([tutorial](#)): `torch.nn.Linear(2048, 1), torch.nn.Sigmoid`

If the class label of Cat is 1, the output value (range: 0 ~ 1)

should be **close to 1** for cat images, and vice versa.

#### (3) Run the function to show the structure **in the terminal**

→ Hint:

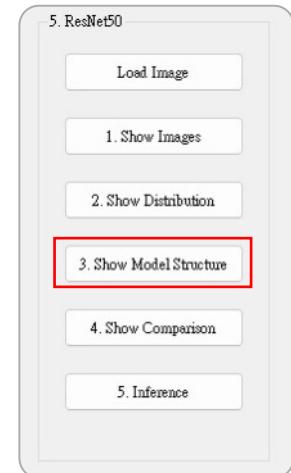
(a) TensorFlow: `tf.keras.Model.summary()`

(b) PyTorch: `torchsummary`

## 2) When the demo:

### (1) Click the button “**3. Show Model Structure**”

### (2) Run the function to show the structure **in the terminal**



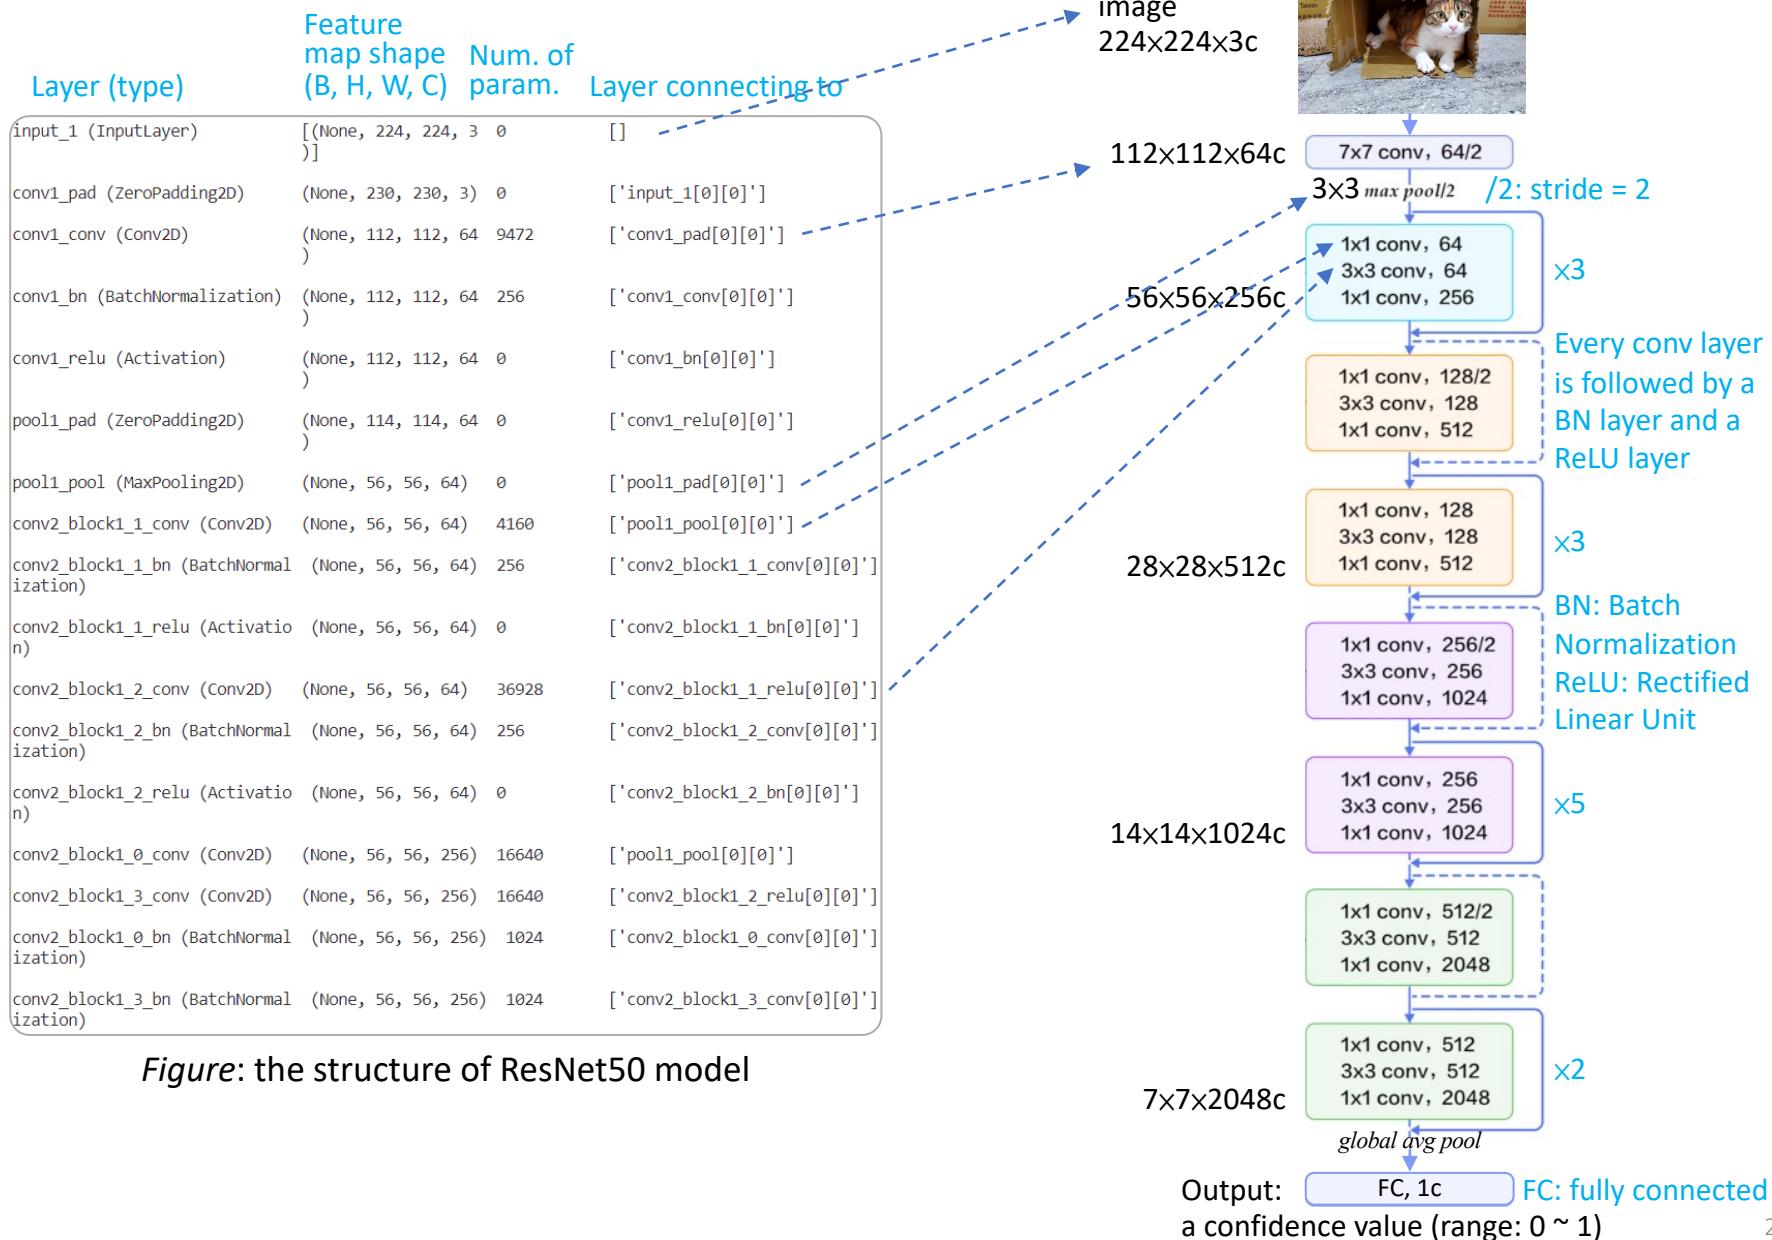
| Layer (type)                            | Feature map shape<br>(B, H, W, C) | Num. of param. | Layer connecting to           |
|-----------------------------------------|-----------------------------------|----------------|-------------------------------|
| input_1 (InputLayer)                    | [None, 224, 224, 3 0 )]           | 0              | []                            |
| conv1_pad (ZeroPadding2D)               | (None, 230, 230, 3) 0             | 0              | ['input_1[0][0]']             |
| conv1_conv (Conv2D)                     | (None, 112, 112, 64 9472 )        | 9472           | ['conv1_pad[0][0]']           |
| conv1_bn (BatchNormalization)           | (None, 112, 112, 64 256 )         | 256            | ['conv1_conv[0][0]']          |
| conv1_relu (Activation)                 | (None, 112, 112, 64 0 )           | 0              | ['conv1_bn[0][0]']            |
| pool1_pad (ZeroPadding2D)               | (None, 114, 114, 64 0 )           | 0              | ['conv1_relu[0][0]']          |
| pool1_pool (MaxPooling2D)               | (None, 56, 56, 64) 0              | 0              | ['pool1_pad[0][0]']           |
| conv2_block1_1_conv (Conv2D)            | (None, 56, 56, 64) 4160           | 4160           | ['pool1_pool[0][0]']          |
| conv2_block1_1_bn (BatchNormaliz ation) | (None, 56, 56, 64) 256            | 256            | ['conv2_block1_1_conv[0][0]'] |
| conv2_block1_1_relu (Activatio n)       | (None, 56, 56, 64) 0              | 0              | ['conv2_block1_1_bn[0][0]']   |

Figure: the structure of ResNet50 model

# 5.3 (3%) Show the structure of ResNet50 model

(出題 : Benjamin)

## 2. How to read the structure



# 5.4 (3%) Set up 2 kinds of loss functions to train 2 ResNet50 models

(出題 : Benjamin)

## 1. Objective

### 1) At home:

- (1) Set up **Focal Loss ( $\alpha$ -balanced)** and **Binary Cross Entropy** in codes for model training (train.py)

→ Hint:

(a) TensorFlow: `tfa.losses.SigmoidFocalCrossEntropy()`, `tf.keras.losses.BinaryCrossEntropy()`

(b) PyTorch: `torchvision.ops.sigmoid_focal_loss()`, `torch.nn.function.binary_cross_entropy()`

Because you have to use **Focal Loss**, which accepts 1 confidence value, you may replace the output layer of ResNet50 model to a FC layer of **1 node** with a **Sigmoid** activation function.

- (2) Train **2 ResNet50 models** with training dataset using **2 loss functions**

→ Hint:

(a) TensorFlow: `tf.keras.Model.fit()`

(b) PyTorch ([tutorial](#)): write a for loop to train the model

### 2) When the demo:

- (1) Show your **codes** about loss functions in train.py

```
60  ### Q5.4 Set up 2 kinds of loss functions to train 2 ResNet50 models
61  # 1st loss function
62  loss_function = tfa.losses.SigmoidFocalCrossEntropy(alpha=0.4, gamma=1.0)
63
64  # 2nd loss function
65  # loss_function = tf.keras.losses.BinaryCrossentropy()
66
67  optimizer = tf.keras.optimizers.Adam(learning_rate=8e-5)
68
69  # configure loss function and optimizer for training
70  model.compile(optimizer=optimizer, loss=loss_function, metrics=['accuracy'])
```

Train a model with Focal Loss

Train another model with Binary Cross Entropy

Figure: Code example (you can follow the hyperparameter settings)

# 5.4 (3%) Set up 2 kinds of loss functions to train 2 ResNet50 models

(出題 : Benjamin)

## 2. What is Focal Loss?

- $FL(p_t) = -\alpha_t(1 - p_t)^\gamma \log(p_t)$

1)  $p_t = \begin{cases} p & \text{if } y = 1 \\ 1 - p & \text{otherwise,} \end{cases}$

$p$ : model's **output confidence** value for the input with class label  $y = 1$  (range:  $0 \sim 1$ )

$y$ : **ground-truth** class label (0 or 1)

2)  $\alpha_t = \begin{cases} \alpha & \text{if } y = 1 \\ 1 - \alpha & \text{otherwise,} \end{cases}$

$\alpha$ : class imbalance weighting factor (Usually set to  $1 - \text{frequency of } y = 1 \text{ examples}$ )

- 3)  $\gamma$ : focusing parameter (Usually set to 2)

- Focal Loss mitigates **class imbalance** by assigning **more weights** to hard examples and less weights to easy examples.

- Focal Loss is for **binary classification**, ex: Foreground and Background, Face and Non-Face, or Cat and Dog.

- Reference: [Focal Loss for Dense Object Detection](#)

Ex: we have

1) 30 examples with  $y = 0$

2) 70 examples with  $y = 1$

→ Frequency of  $y = 1$  examples is  $\frac{70}{100} = 0.7$

→  $\alpha = 1 - 0.7 = 0.3$

# 5.5 (4%) Compare the accuracies of 2 ResNet50 models on validation dataset

(出題 : Benjamin)

## 1. Objective

### 1) At home:

- (1) Validate **2 ResNet50 models** with validation dataset  
→ Hint:
  - (a) TensorFlow: `tf.keras.Model.evaluate()`
  - (b) PyTorch ([tutorial](#)): write a for loop to validate the model

- (2) Plot the **accuracy values** with a **bar chart**

- (3) Save the figure

### 2) When the demo:

- (1) Click the button "**4. Show Comparison**"
- (2) Show the **saved figure** of accuracy comparison in a new window

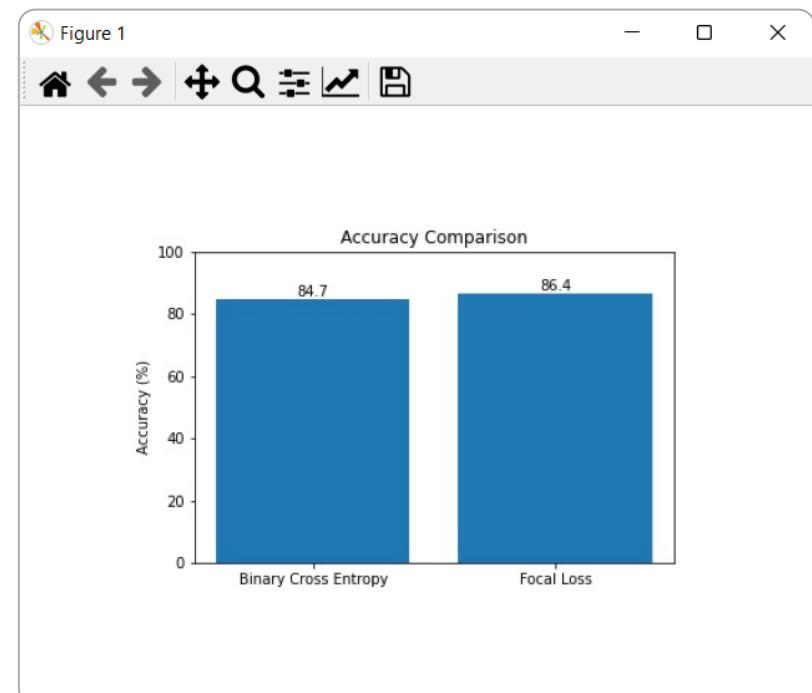
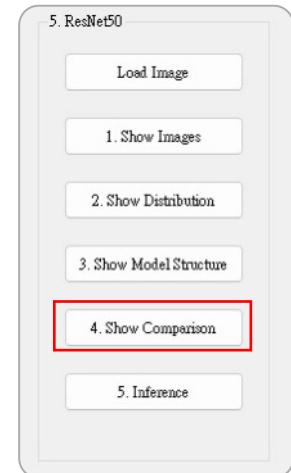


Figure: Accuracy Comparison

Notice: this is an example, the numbers might differ

# 5.6 (4%) Use the better-trained model to run inference and show the predicted class label

(出題 : Benjamin)

## 1. Objective

### 1) At home:

#### (1) Load the trained model

→ Hint:

- (a) TensorFlow: `tf.keras.Model.load_weights()`
- (b) PyTorch: `torch.nn.Module.load_state_dict()`

#### (2) Click the button “Load Image” to select 1 image arbitrarily

→ Hint: `PyQt5.QtWidgets.QFileDialog.getOpenFileName()`

#### (3) Show the loaded image in the GUI

#### (4) Resize the loaded image to 224x224x3c (RGB)

#### (5) Click the button “5. Inference” to run inference on the resized image

→ Hint:

- (a) TensorFlow: `tf.keras.Model.predict_step()`
- (b) PyTorch: pass an image when calling `torch.nn.Module.inference`, ex: `trained_model(img)`

#### (6) Show the predicted class label

→ Hint: decide the class label with a **threshold** of the **output value**.

Ex: class label =  $\begin{cases} \text{Cat}, & \text{output} < \text{thresh} \\ \text{Dog}, & \text{output} \geq \text{thresh} \end{cases}$   
 $\text{thresh} = 0.5$

### 2) When the demo: **repeat** the process

